

```
pip install statsmodels
```

```
pip install scipy
```

```
Requirement already satisfied: scipy in c:\users\nikhilesh\jupyter_env\lib\site-packages (1.16.1)
```

```
Requirement already satisfied: numpy<2.6,>=1.25.2 in c:\users\nikhilesh\jupyter_env\lib\site-packages (from scipy) (2.3.2)
```

```
Note: you may need to restart the kernel to use updated packages.
```

```
[notice] A new release of pip is available: 25.1.1 -> 25.2
```

```
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("Churn_Modelling.csv")
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

## Mean

```
mean = np.mean(df["Age"])
print(mean)

38.9218
```

## Check Data Type of Columns

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore            10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard               10000 non-null  int64
11  IsActiveMember          10000 non-null  int64
12  EstimatedSalary         10000 non-null  float64
13  Exited                  10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

## Mean Absolute Deviation

```
mad = np.sum(np.abs(df["Age"] - mean))/len(df["Age"])
print(mad)

7.94097904

print(mean)

38.9218
```

## Check Standard Deviation


```
np.std(df["Age"]) , np.var(df["Age"])  
(np.float64(10.487282048271611), np.float64(109.98308476))
```

## Interquartile Range

```
Q1 = df["Age"].quantile(0.25)  
Q3 = df["Age"].quantile(0.75)  
IQR = Q3 - Q1  
  
lower_fence = Q1 - 1.5 * IQR  
upper_fence = Q3 + 1.5 * IQR  
  
outliers = df[(df["Age"] < lower_fence) | (df["Age"] > upper_fence)]  
  
print("Q1:", Q1, "Q3:", Q3, "IQR:", IQR)  
print("Outliers count:", outliers.shape[0])  
  
Q1: 32.0 Q3: 44.0 IQR: 12.0  
Outliers count: 359
```

## Interquartile Range with plot

```
sns.boxplot(x="Age", data = df)  
plt.show()
```



## Check mean , STD , Minimun , Maximum

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age
Tenure \				
count	10000.000000	1.000000e+04	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800
std	2886.89568	7.193619e+04	96.653299	10.487806
min	1.00000	1.556570e+07	350.000000	18.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000
max	10000.00000	1.581569e+07	850.000000	92.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
mean	76485.889288	1.530200	0.70550	0.515100	
std	62397.405202	0.581654	0.45584	0.499797	
min	0.000000	1.000000	0.00000	0.000000	
25%	0.000000	1.000000	0.00000	0.000000	
50%	97198.540000	1.000000	1.00000	1.000000	
75%	127644.240000	2.000000	1.00000	1.000000	
max	250898.090000	4.000000	1.00000	1.000000	

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

## Check Skewness


```
from scipy.stats import skew, kurtosis

print("Skewness:", skew(df["Age"].dropna()))
print("Kurtosis:", kurtosis(df["Age"].dropna()))
```


Skewness: 1.0111685586628076  
Kurtosis: 1.3940495456392599

## Skewness plot


```
sns.histplot(x= "Age" , data= df)
plt.show()
```



```
plt.figure(figsize=(6,4))
sns.countplot(x = "Age" , data = df)
plt.show()
```



```
sns.histplot(df["Age"].dropna(), kde = True)
plt.show()
```




```
survival_prob = df['HasCrCard'].value_counts(normalize=True)
```

```
plt.figure(figsize=(6,4))
```

```
sns.barplot(x=survival_prob.index, y=survival_prob.values)
```

```
<Axes: xlabel='HasCrCard'>
```



```
df.tail()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
9995	9996	15606229	Obijiaku	771	France	Male
39						
9996	9997	15569892	Johnstone	516	France	Male
35						
9997	9998	15584532	Liu	709	France	Female
36						
9998	9999	15682355	Sabbatini	772	Germany	Male
42						
9999	10000	15628319	Walker	792	France	Female
28						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
9995	5	0.00	2	1		0
9996	10	57369.61	1	1		1
9997	7	0.00	1	0		1
9998	3	75075.31	2	1		0
9999	4	130142.79	1	1		0



	EstimatedSalary	Exited
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

## Check Correlation


```
corr = df.corr(numeric_only=True)
print(corr['Age'].sort_values(ascending=False))

# Heatmap visualization
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f")
```

Age	1.000000
Exited	0.285323
IsActiveMember	0.085472
Balance	0.028308
CustomerId	0.009497
RowNumber	0.000783
CreditScore	-0.003965
EstimatedSalary	-0.007201
Tenure	-0.009997
HasCrCard	-0.011721
NumOfProducts	-0.030680

Name: Age, dtype: float64

<Axes: >



## Check P-Vlaue

```
from scipy.stats import chi2_contingency

table = pd.crosstab(df['HasCrCard'], df['Exited'])
print("Contingency Table:\n", table)

# Chi-square test
chi2, p, dof, expected = chi2_contingency(table)

print("\nChi-square Statistic:", chi2)
print("Degrees of Freedom:", dof)
print("p-value:", p)
```

Contingency Table:

	Exited	0	1
HasCrCard	0	2332	613


```
1          5631  1424
```

```
Chi-square Statistic: 0.47133779904440803
```

```
Degrees of Freedom: 1
```

```
p-value: 0.49237236141554697
```

```
sns.histplot(df["Age"], bins=30, kde=True, color="skyblue")  
plt.show()
```



```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
--------	---------	---------------	-----------	----------------	---

0	2	0.00	1	1	1
1	1	83807.86	1	0	1
2	8	159660.80	3	1	0
3	1	0.00	2	0	0
4	2	125510.82	1	1	1

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

## Check Normal Distribution

```
from scipy.stats import normaltest

stat, p = normaltest(df["Age"])
print("Test Statistic:", stat)
print("p-value:", p)

if p < 0.05:
    print("Conclusion: Age distribution is NOT perfectly normal.")
else:
    print("Conclusion: Age distribution follows normal distribution.")
```

Test Statistic: 1507.7908881363314  
p-value: 0.0  
Conclusion: Age distribution is NOT perfectly normal.

## Check Covariance

```
num_df = df[['Age', 'Tenure']].dropna()

# Covariance matrix
cov_matrix = num_df.cov()
print("Covariance Matrix:\n", cov_matrix)


# Individual covariance value
cov_age_fare = num_df['Age'].cov(num_df['Tenure'])
print("\nCovariance between Age and Fare:", cov_age_fare)
```

Covariance Matrix:

	Age	Tenure
Age	109.994084	-0.303229
Tenure	-0.303229	8.364673

Covariance between Age and Fare: -0.30322936293629343

```
num_df = df[['Age', 'Tenure']].dropna()
sns.scatterplot(x='Age', y='Tenure', data=num_df)
plt.show()
```



## Central Limit Theorem

```
sample_means = []
for i in range(1000):
    sample = np.random.choice(df["Age"], size=80, replace=True)
    sample_means.append(sample.mean())


# Plot original age distribution
plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
sns.histplot(df["Age"], bins=30, kde=True, color="skyblue")
plt.title("Original Age Distribution (Not Normal)")

# Plot sampling distribution of sample means
plt.subplot(1,2,2)
```

```
sns.histplot(sample_means, bins=30, kde=True, color="orange")
plt.title("Distribution of Sample Means (Approx Normal)")

Text(0.5, 1.0, 'Distribution of Sample Means (Approx Normal)')
```



```
np.mean(sample_means)
np.float64(38.883925)

np.mean(df["Age"])
np.float64(38.9218)
```

## Z- Test

```
from statsmodels.stats.weightstats import ztest

data = df['Age']

# Perform one-sample Z-test (test mean = 5.8)
z_stat, p_val = ztest(data, value=5.8)

print("Z-statistic:", z_stat)
print("p-value:", p_val)

# Decision rule
alpha = 0.05
if p_val < alpha:
    print("Reject Null Hypothesis → Mean is significantly different
    from 5.8")
```

```

else:
    print("Fail to Reject Null → No significant difference from 5.8")

Z-statistic: 315.81246424142995
p-value: 0.0
Reject Null Hypothesis → Mean is significantly different from 5.8

data = df['Age']

# Hypothesized population mean
mu = 38.92

# Perform one-sample Z-test
z_stat, p_val = ztest(data, value=mu)


# Plot histogram + KDE
sns.histplot(data, bins=20, kde=True, color="skyblue", stat="density")

# Plot hypothesized mean line
plt.axvline(mu, color='red', linestyle='--', linewidth=2,
label=f"Hypothesized Mean = {mu}")

# Plot sample mean line
sample_mean = np.mean(data)
plt.axvline(sample_mean, color='green', linestyle='-', linewidth=2,
label=f"Sample Mean = {sample_mean:.2f}")

<matplotlib.lines.Line2D at 0x1c87bc49d10>

```



## Z-Test with t\_table and t\_calculate

```
import scipy.stats as st

z_t = st.norm.ppf(0.95)
print(z_t)

1.6448536269514722

z_score = (10.487282048271611 / np.sqrt(1000))
print(z_score)

0.331636977371342

if z_table < z_score:
    print("h1 is correct")
else:
    print("h0 is correct")

h0 is correct
```



## Check Mean with Sample Data

```
sample_mean = df.loc[:4000, "Age"].mean()
print(sample_mean)

38.895526118470386
```

## Z-test with t\_table and t\_Calculate

```
import scipy.stats as st

z_table = st.norm.ppf(0.95)
print(z_table)

1.6448536269514722

std = np.std(df["Age"])
print(std)

10.487282048271611

z_cal = (38.89 - 38.92)/(10.48/np.sqrt(4000))
print(z_cal)

-0.18104643092567743
```

## Z-test with Plot

```
from statsmodels.stats.weightstats import ztest

data = df['Age']

# Hypothesized population mean
mu = 38.92

# Perform one-sample Z-test
z_stat, p_val = ztest(data, value=mu)

# Plot histogram + KDE
plt.figure(figsize=(8,5))
sns.histplot(data, bins=80, kde=True, color="skyblue", stat="density")

# Plot hypothesized mean line
plt.axvline(mu, color='red', linestyle='--', linewidth=2,
label=f"Hypothesized Mean = {mu}")

# Plot sample mean line
#sample_mean = np.mean(data)
```


```

sample_mean = df.loc[:4000, "Age"].mean()
plt.axvline(sample_mean, color='green', linestyle='--', linewidth=2,
label=f"Sample Mean = {sample_mean:.2f}")

# Title and labels
plt.title("Z-Test Visualization on Iris Sepal Length")
plt.xlabel("Sepal Length (cm)")
plt.ylabel("Density")
plt.legend()
plt.show()

print("Z-statistic:", z_stat)
print("p-value:", p_val)

```



```

Z-statistic: 0.017162788122422362
p-value: 0.9863067485904023

```

## T-Test

```

from scipy.stats import ttest_ind

df = df.dropna(subset=["Exited", "HasCrCard"])

```

```

# Split male and female groups
exited = df[df["HasCrCard"] == 0]["Exited"]
nonexited = df[df["HasCrCard"] == 1]["Exited"]

# Perform independent T-test
t_stat, p_val = ttest_ind(exited, nonexited)

print("T-statistic:", t_stat)
print("p-value:", p_val)

# Decision rule
alpha = 0.05
if p_val < 0.05:
    print("Reject Null Hypothesis.")
else:
    print("Fail to Reject Null")

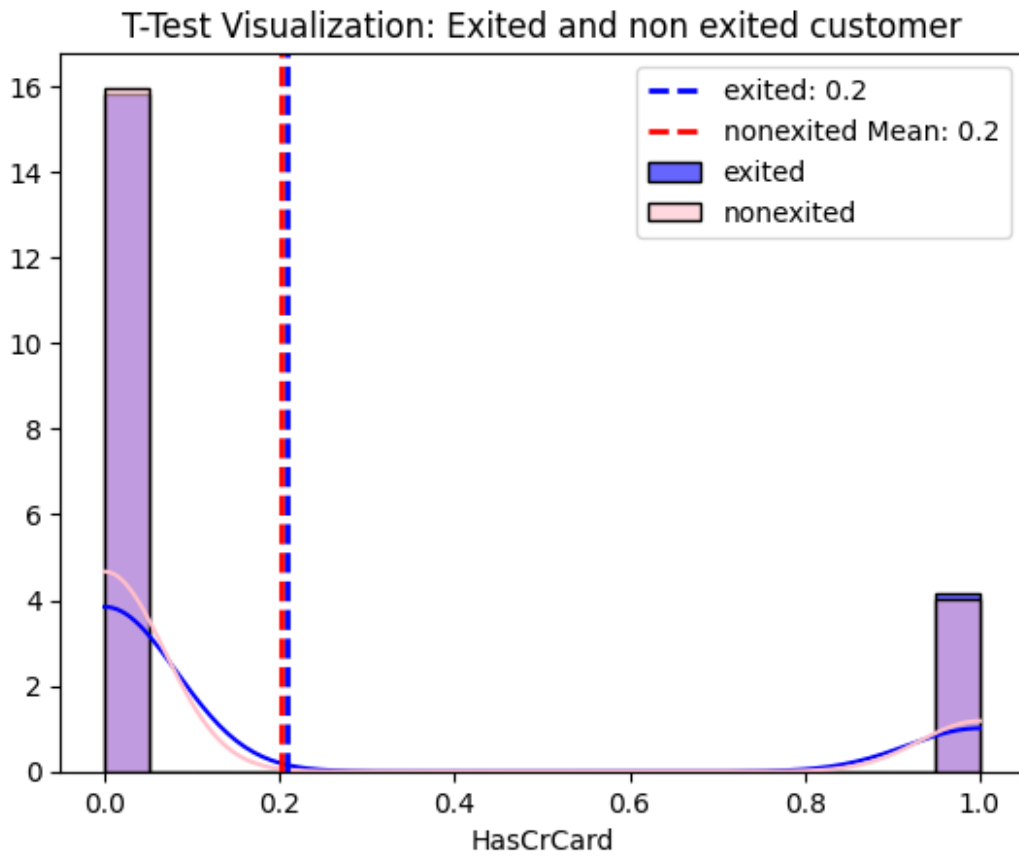
T-statistic: 0.7137233605912553
p-value: 0.47541491837605965
Fail to Reject Null

sns.histplot(exited, bins=20, kde=True, color="blue", label="exited",
stat="density", alpha=0.6)
sns.histplot(nonexited, bins=20, kde=True, color="pink",
label="nonexited", stat="density", alpha=0.6)

plt.axvline(exited.mean(), color="blue", linestyle="--", linewidth=2,
label=f"exited: {male_mass.mean():.1f}")
plt.axvline(nonexited.mean(), color="red", linestyle="--",
linewidth=2, label=f"nonexited Mean: {female_mass.mean():.1f}")

plt.title("T-Test Visualization: Exited and non exited customer")
plt.xlabel("HasCrCard")
plt.ylabel("")
plt.legend()
plt.show()

```



## Check Causation

```
import statsmodels.api as sm
df.dropna()

# Example: Does total_bill cause tip amount?
X = df["HasCrCard"]      # Independent variable
y = df["Exited"]         # Dependent variable

# Add constant for intercept
X = sm.add_constant(X)

# Fit linear regression model
model = sm.OLS(y, X).fit()

# Summary gives p-values and R-squared (to infer causation with caution)
print(model.summary())
```

OLS Regression Results

=====

```

=====
Dep. Variable:                Exited    R-squared:
0.000
Model:                        OLS      Adj. R-squared:
-0.000
Method:                       Least Squares    F-statistic:
0.5094
Date:                         Mon, 15 Sep 2025    Prob (F-statistic):
0.475
Time:                         18:02:07    Log-Likelihood:
-5094.7
No. Observations:             10000    AIC:
1.019e+04
Df Residuals:                 9998    BIC:
1.021e+04
Df Model:                     1

Covariance Type:              nonrobust

=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const                0.2081      0.007      28.045      0.000      0.194
0.223
HasCrCard            -0.0063      0.009      -0.714      0.475      -0.024
0.011
=====
=====
Omnibus:              2043.753    Durbin-Watson:
1.994
Prob(Omnibus):        0.000    Jarque-Bera (JB):
3619.106
Skew:                 1.471    Prob(JB):
0.00
Kurtosis:             3.165    Cond. No.
3.45
=====
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.

```

1. Dep. Variable: Exited

Ye dependent variable hai (target) — matlab hum predict kar rahe hain ki customer Exited (churn kiya) ya nahi.

1. R-squared: 0.000

$R^2$  measure karta hai ki model dependent variable ka kitna variance explain karta hai.

0.000 matlab model almost kuch bhi explain nahi kar raha (predictive power  $\approx$  zero).

1. F-statistic: 0.5094, Prob (F-statistic): 0.475

F-test check karta hai ki model overall significant hai ya nahi.

p-value (0.475)  $>$  0.05  $\rightarrow$  model significant nahi hai  $\rightarrow$  explanatory variable meaningful nahi hai.

1. coef (coefficients)

const = 0.2081  $\rightarrow$  Intercept (baseline probability of exit  $\sim$ 20.8%).

HasCrCard = -0.0063  $\rightarrow$  Agar customer ke paas credit card hai to exit probability  $\sim$ 0.6% kam hoti hai.

1.  $P > |t|$  values (Hypothesis testing)

For HasCrCard,  $p = 0.475$  ( $>$  0.05). Matlab HasCrCard ka effect statistically significant nahi hai.

Hum null hypothesis (no effect) reject nahi kar paate.

1. Confidence Interval [0.025, 0.975]

HasCrCard ka interval = [-0.024, 0.011].

Zero is inside interval  $\rightarrow$  again, effect not significant.

1. Other stats

Durbin-Watson = 1.994  $\rightarrow$  Residuals ka autocorrelation theek hai ( $\sim$ 2 = good).

Omnibus / Jarque-Bera  $\rightarrow$  Residuals normality test (significant  $\rightarrow$  normality issue).

Summary (Simple Words)

Model explain nahi kar raha ( $R^2 = 0$ ).

Predictor (HasCrCard) ka effect statistically insignificant hai ( $p = 0.475$ ).

Matlab credit card hone ka churn par koi meaningful impact nahi hai is dataset me.

"OLS results show ki HasCrCard ka coefficient -0.0063 hai but p-value 0.475 hai ( $>$ 0.05). Iska matlab hai ki credit card hone ka customer churn par koi statistically significant impact nahi hai.  $R^2$  value bhi  $\sim$ 0 hai, jo dikhata hai ki model exit variable ko explain nahi kar pa raha."

```
# Probabli
```

```
total = len(df)
```

```
# Example 1: Probability of survival
```

```
p_exited = df['Exited'].sum() / total
```

```

# Example 2: Probability of male passenger
p_male = (df['Gender'] == 'Male').sum() / total

# Example 3: Joint probability → male AND survived
p_male_exited = len(df[(df['Gender']=='Male') & (df['Exited']==1)]) /
total

# Example 4: Conditional probability → P(Survived | Female)
p_exited_given_female = len(df[(df['Gender']=='Female') &
(df['Exited']==1)]) / (df['Gender']=='Female').sum()

print("P(Exited):", round(p_exited, 3))
print("P(Male):", round(p_male, 3))
print("P(Male n Exited):", round(p_male_exited, 3))
print("P(Exited | Female):", round(p_exited_given_female, 3))

P(Exited): 0.204
P(Male): 0.546
P(Male n Exited): 0.09
P(Exited | Female): 0.251

```